

## Aim: Linear Regression by using Deep Neural Network

```
In [ ]: pip install tensorflow --user --no-warn-script-location
```

```
In [1]: import tensorflow as tf
        from tensorflow.keras.datasets import boston_housing
        from sklearn import preprocessing
```

```
In [2]: (train_x, train_y),(test_x,test_y)=boston_housing.load_data()
```

```
In [3]: print(train_x.shape)
        print(test_x.shape)
        print(train_y.shape)
        print(test_y.shape)
```

```
(404, 13)
(102, 13)
(404,)
(102,)
```

```
In [4]: train_x[0]
```

```
Out[4]: array([ 1.23247,  0.      ,  8.14   ,  0.      ,  0.538   ,  6.142   ,
                91.7    ,  3.9769 ,  4.     , 307.    ,  21.     , 396.9    ,
                18.72   ])
```

```
In [5]: train_y[0]
```

```
Out[5]: 15.2
```

```
In [6]: train_x=preprocessing.normalize(train_x)
        test_x=preprocessing.normalize(test_x)
```

```
In [7]: train_x[0]
```

```
Out[7]: array([0.0024119 , 0.          , 0.01592969, 0.          , 0.00105285,
                0.01201967, 0.17945359, 0.00778265, 0.00782786, 0.6007879 ,
                0.04109624, 0.77671895, 0.03663436])
```

```
In [8]: train_y[0]
```

```
Out[8]: 15.2
```

```
In [9]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import *

        def HPPM():
            model=Sequential()
            model.add(Dense(128, activation='relu', input_shape=(train_x[0].shape)))
            model.add(Dense(64, activation='relu'))
            model.add(Dense(32, activation='relu'))
            model.add(Dense(1))
            model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
            return model
```

```
In [10]: import numpy as np
         k=4
         num_val_samples=len(train_x)
         num_epochs=100
         all_scores=[]
```

```
In [11]: model=HPPM()
         history=model.fit(x=train_x, y=train_y, epochs=num_epochs, batch_size=1, verbose=1)
```

```
Epoch 1/100
404/404 [=====] - 1s 2ms/step - loss: 145.2648 -
mae: 8.8432 - val_loss: 70.4067 - val_mae: 6.0675
Epoch 2/100
404/404 [=====] - 2s 4ms/step - loss: 69.6452 -
mae: 5.8083 - val_loss: 60.2890 - val_mae: 5.7936
Epoch 3/100
404/404 [=====] - 2s 4ms/step - loss: 64.0329 -
mae: 5.4785 - val_loss: 58.2462 - val_mae: 5.5318
Epoch 4/100
404/404 [=====] - 2s 4ms/step - loss: 62.4070 -
mae: 5.4575 - val_loss: 55.6510 - val_mae: 5.6276
Epoch 5/100
404/404 [=====] - 1s 4ms/step - loss: 57.9584 -
mae: 5.3383 - val_loss: 54.0462 - val_mae: 5.4752
Epoch 6/100
404/404 [=====] - 1s 3ms/step - loss: 58.4756 -
mae: 5.2441 - val_loss: 61.6325 - val_mae: 5.5175
Epoch 7/100
404/404 [=====] - 1s 3ms/step - loss: 55.3075 -
mae: 5.1750 - val_loss: 55.3075 - val_mae: 5.1750
```

```
In [12]: test_input=[(8.65407330e-05, 0.00000000e+00, 1.13392175e-02, 0.00000000e+00,
                    print("Actual output : 15.2")
                    print("Predicted output: ", model.predict(test_input))
```

```
Actual output : 15.2
1/1 [=====] - 0s 105ms/step
Predicted output: [[15.099294]]
```

