

## Aim: Convolutional Neural Network (CNN)

```
In [1]: import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow import keras
import numpy as np

(x_train, y_train), (x_test, y_test) = keras.datasets.fashion_mnist.load_data
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz (https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz)
29515/29515 [=====] - 0s 2us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz (https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz)
26421880/26421880 [=====] - 176s 7us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz (https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz)
5148/5148 [=====] - 0s 0s/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz (https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz)
4422102/4422102 [=====] - 8s 2us/step
```

```
In [2]: x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
In [3]: x_train.shape
```

```
Out[3]: (60000, 28, 28, 1)
```

```
In [4]: x_test.shape
```

```
Out[4]: (10000, 28, 28, 1)
```

```
In [5]: y_train.shape
```

```
Out[5]: (60000,)
```

```
In [6]: y_test.shape
```

```
Out[6]: (10000,)
```

```
In [7]: model = keras.Sequential([
    keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),

    keras.layers.MaxPooling2D((2,2)),

    keras.layers.Dropout(0.25),

    keras.layers.Conv2D(64, (3,3), activation='relu'),

    keras.layers.MaxPooling2D((2,2)),

    keras.layers.Dropout(0.25),

    keras.layers.Conv2D(128, (3,3), activation='relu'),

    keras.layers.Flatten(),
    keras.layers.Dense(128, activation='relu'),

    keras.layers.Dropout(0.25),
    keras.layers.Dense(10, activation='softmax')
])
```

```
In [8]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
dropout_1 (Dropout)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 128)	73856
flatten (Flatten)	(None, 1152)	0
dense (Dense)	(None, 128)	147584
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
=====		
Total params: 241,546		
Trainable params: 241,546		
Non-trainable params: 0		

```
In [10]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=3, validation_data=(x_test, y_test))
```

```
Epoch 1/3
1875/1875 [=====] - 46s 24ms/step - loss: 0.4153 - accuracy: 0.8474 - val_loss: 0.3430 - val_accuracy: 0.8802
Epoch 2/3
1875/1875 [=====] - 46s 24ms/step - loss: 0.3342 - accuracy: 0.8773 - val_loss: 0.3052 - val_accuracy: 0.8883
Epoch 3/3
1875/1875 [=====] - 45s 24ms/step - loss: 0.3038 - accuracy: 0.8875 - val_loss: 0.2720 - val_accuracy: 0.8999
```

```
In [11]: test_loss, test_acc = model.evaluate(x_test, y_test)
```

```
print('Test accuracy:', test_acc)
```

```
313/313 [=====] - 2s 7ms/step - loss: 0.2720 - accu  
racy: 0.8999
```

```
Test accuracy: 0.8999000191688538
```